

CS 630, Fall 2024, Homework 6

Yifei Bao

November 11, 2024

Problem 1 *Bloom Filter Transformations*

1. In Bloom filter B , each element x has r hash functions that map the element to specific positions in the bit array of B , setting those bit positions to 1.

In B' , the original bit array is split into B_1 and B_2 , and B' is formed by taking the bit-wise OR of these two halves, $B' = B_1 \vee B_2$.

The new hash values in B' are computed as $h_i(x) \bmod (m/2)$. This means that any bit that was set in either B_1 or B_2 will be folded as a 1 in B' because of boolean or. Therefore, for every element originally in B , the corresponding bits will still be set in B' , and the lookup will always return positive.

2.

a. For items that were in B prior to creating B' , False negative is when an element that is in the set returns 0 during a lookup. Since B' is created by taking the bit-wise or of B_1 and B_2 , it keeps all the bit positions set by the original elements. Therefore, the false negative rate is 0.

False positive is when an element that is not in the set returns 1 during a lookup. The false positive rate of B is $p = (1 - e^{-rn/m})^r$. The false positive rate of B' is $p' = (1 - e^{-rn/(m/2)})^r$. Since $m/2 < m$, the false positive rate of B' is larger than that of B .

b. For items that were not in B , since they were not in B , the false negative rate is 0 still.

False positive rate is the same as in part a. The false positive rate of B is $p = (1 - e^{-rn/m})^r$. The false positive rate of B' is $p' = (1 - e^{-rn/(m/2)})^r$. Since $m/2 < m$, the false positive rate of B' is larger than that of B .

c. $O(n)$ more items are inserted into B' , the false negative rate is still 0.

The false positive rate of B' is $p' = (1 - e^{-r(n+n)/(m/2)})^r = (1 - e^{-r(2n)/(m/2)})^r$.

Problem 2 *Hash Evaluations for Bloom Filters*

1. When inserting an item into the Bloom filter, the hash functions are invoked $k = \lfloor \ln 2(m/n) \rfloor$ times.

2. When checking an item that was inserted into the Bloom filter, the average number of invocations of the hash functions is $k = \lfloor \ln 2(m/n) \rfloor$. Because you know the item was inserted, you can check all the hash values to see if they are set to 1.

3. The probability that a bit is set to 0 after n items were inserted is $p = (1 - 1/m)^{kn}$. The average number of invocations of the hash functions is $1/p$.

Problem 3 *hash*

1. We need at least two queries because we have two unknowns in a linear equation.

Time complexity: $O(1)$, space complexity: $O(1)$.

Algorithm 1: Calculate()

```

/* Call the hash function  $h(x) = ax + b \pmod{2^k}$  to get  $h(0)$  and  $h(1)$  */
1  $b = h(0)$ ;
2  $temp = h(1)$ ;
3  $a = (temp - b) \pmod{2^{64}}$ ;
4 return a, b

```

2. Since $h(x)$ always results in an odd number, it is obvious that $(ax + b) \pmod{2} = 1$, which is the same as $((a \pmod{2})(x \pmod{2}) + (b \pmod{2})) \pmod{2} = 1$. When x is even, then the equation becomes $(b \pmod{2}) \pmod{2} = 1$, that b is odd. When x is odd, then the equation becomes $(a \pmod{2}) = 0$, that a is even.

To avoid the problem, we can make sure that a is odd, then the hash function will produce both even and odd values.

3. To identify a set of keys that could be chosen to cluster in $1/1024$ of the keys, we can choose $x = n * 2^{k-10}$, thus $h(x) = a(n * 2^{k-10}) + b \pmod{2^k}$. In this way, the keys are clustered in $1/1024$ of the keys.

To avoid the problem, we can select a prime m close to 2^k , and the hash function is $(ax + b) \pmod{m}$. Using a prime modulus breaks the structure that the attacker exploits because powers of two have properties that avoid for such attacks.