

CS630 Graduate Algorithms

October 22, 2024

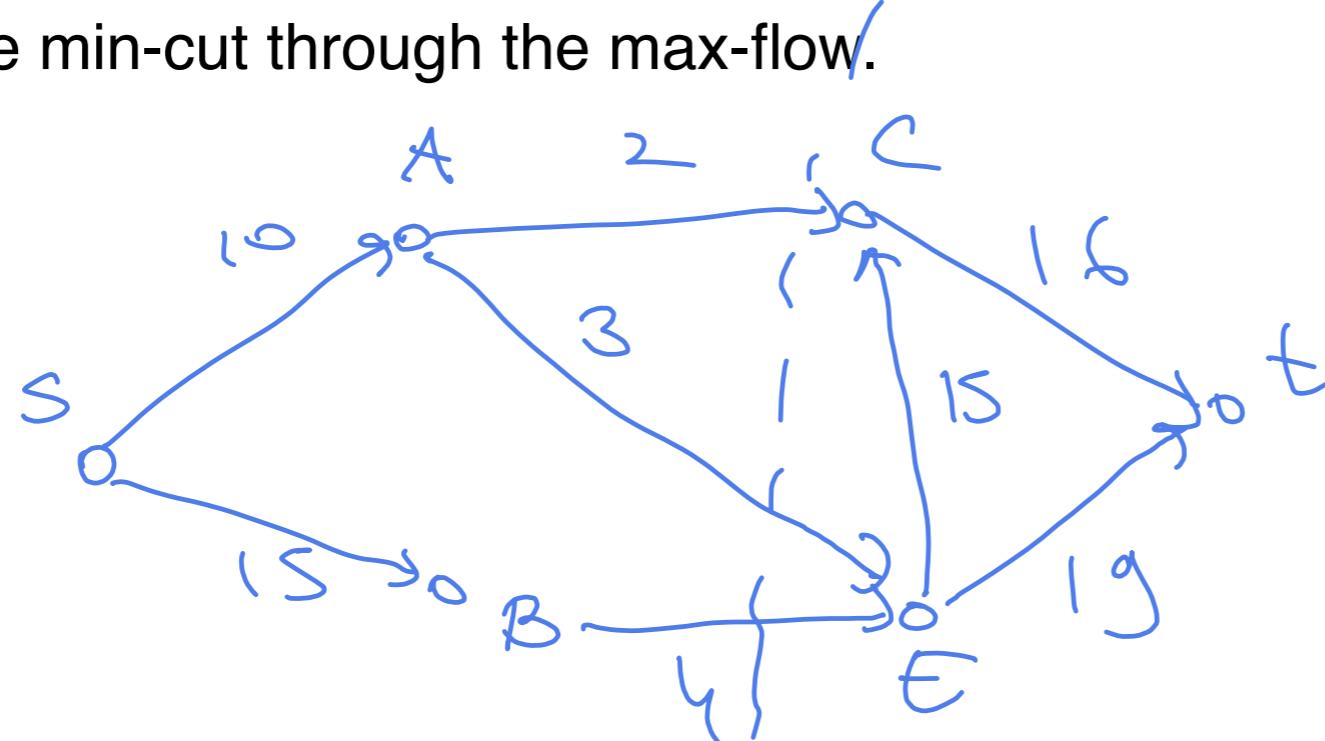
by Dora Erdos and Jeffrey Considine

- min-cut KT ch. 13.2

min st-cut in directed graphs

Max Flow Min Cut theorem: given a directed graph $G(V,E)$ with source s , sink t and edge capacities the value of the max flow is = the capacity of the min st-cut

⇒ find the min-cut through the max-flow.



{s, A, B}

{C, E, T}

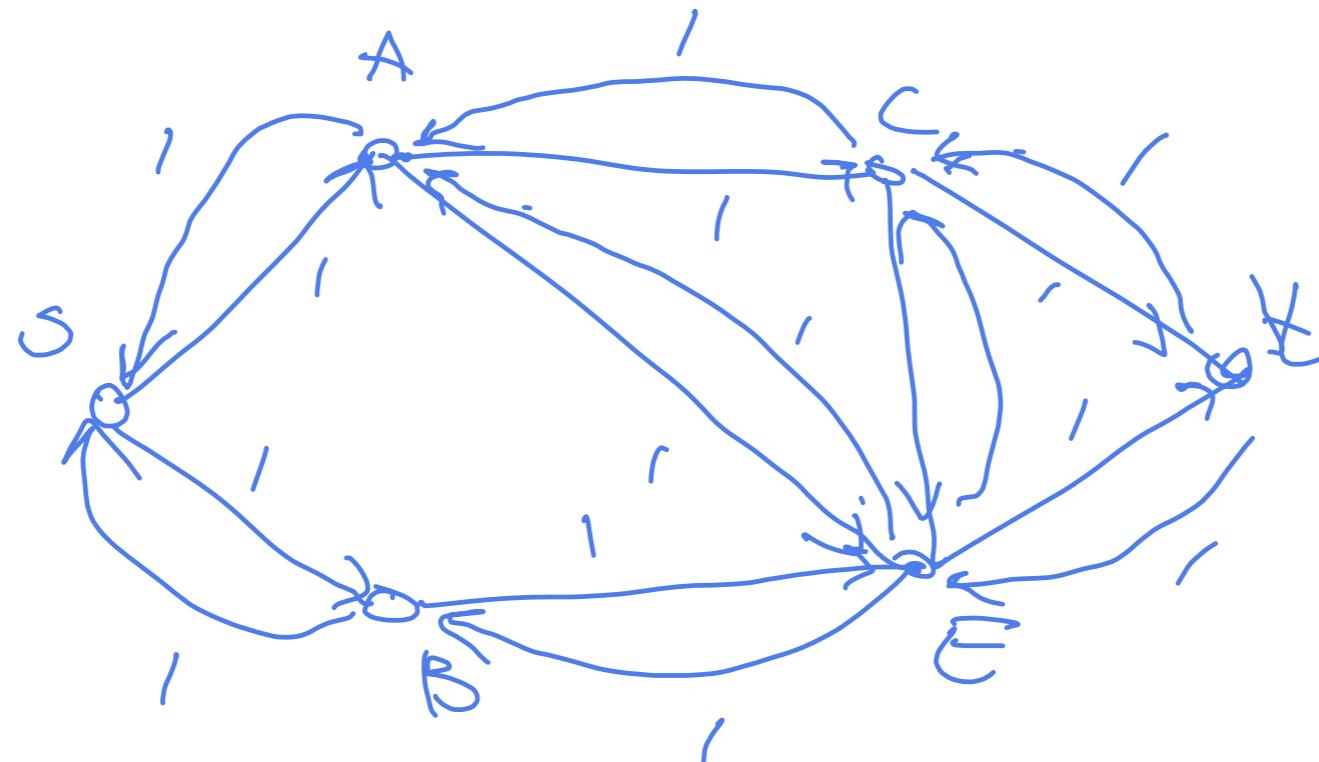
min st - cut: separate s and t while the lowest total capacity between the two sets

min st-cut in undirected unweighted graph

Find the min st-cut in the undirected, unweighted graph $G(V,E)$ with source s and sink t.

capacity/weight of a cut = number of edges between S and V-S.

How?



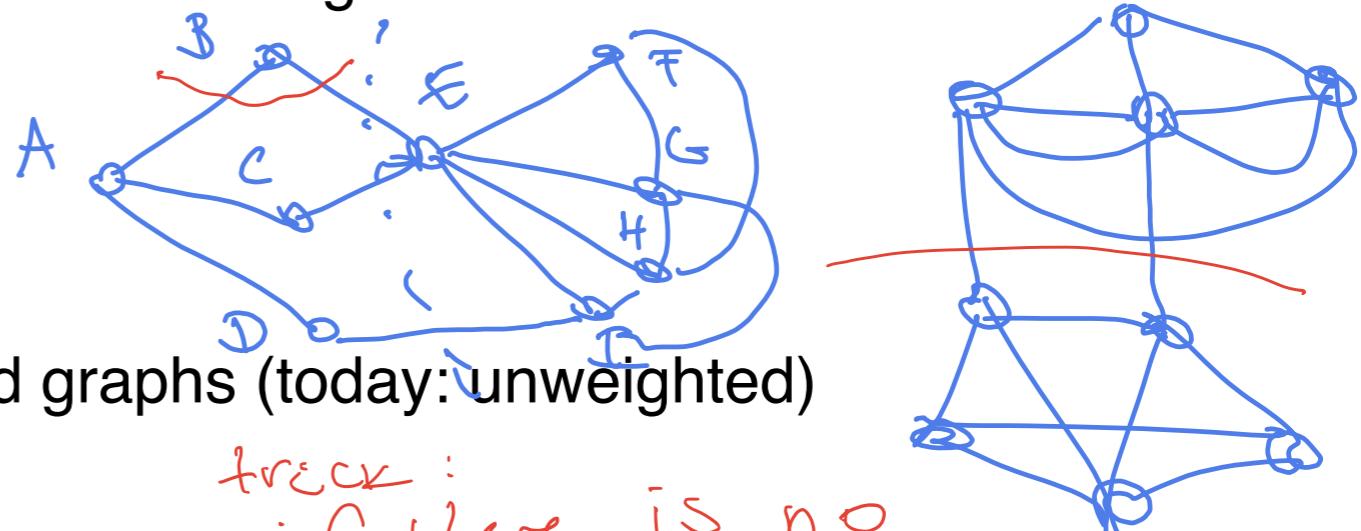
find the
min st-cut
here

$$\binom{n}{2} = \frac{n(n-1)}{2}$$

Global min-cut

Global min-cut: Given an undirected graph $G(V, E)$, find a cut that partitions the nodes into two sets A and $V-A$ with minimum weight.

- no designated source or sink
- $\text{weight}(C) = \sum_{\text{edge } e \text{ in } C} w(e)$
- works for unweighted and weighted graphs (today: unweighted)



Deterministic algorithm (using st-cuts):

For every pair of nodes u, v :

- pick two nodes u and v to serve as source and sink
- find the corresponding min-cut
- keep track of the lowest min-cut so far

runtime: $\binom{n}{2} = O(n^2)$ pairs $\rightarrow O(n^2)$ calls to Ford-Fulkerson

Global min-cut

Global min-cut: Given an undirected graph $G(V,E)$, find a cut that partitions the nodes into two sets A and $V-A$ with minimum weight.

- no designated source or sink
- $\text{weight}(C) = \sum_{\text{edge } e \text{ in } C} w(e)$
- works for unweighted and weighted graphs (today: unweighted)
non-random
II

Deterministic algorithm:

For each pair of vertices u, v assign u as source v as sink, and run the min st-cut algorithm

- since this is an undirected graph, it doesn't matter which of u and v is the source

running time:

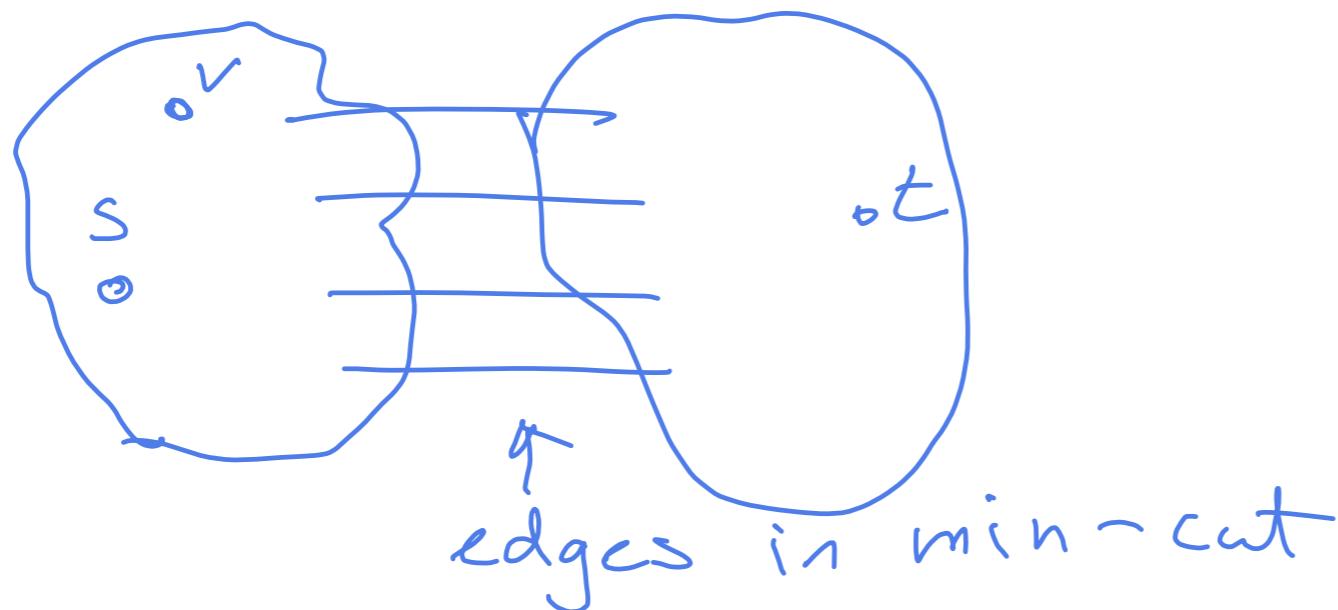
- there are $O(n^2)$ iterations
- each requires a run of FF. We know that C is at most the max degree D .
- in total $O(n^2mD)$

Global min-cut

Global min-cut: Given an undirected graph $G(V,E)$, find a cut that partitions the nodes into two sets A and $V-A$ with minimum weight.

Deterministic algorithm – speed up:

claim: it's enough to fix one source s , and compute the min-cut between s and every other node. $\rightarrow n \cdot \text{calls to FF.}$



vis on the s -side
of the min-cut.

1. there is no min-cut separating s from v
2. any min-cut from v to t is the same as s to t .

runtime:

Global min-cut

Global min-cut: Given an undirected graph $G(V,E)$, find a cut that partitions the nodes into two sets A and $V-A$ with minimum weight.

Deterministic algorithm — speed up:

claim: it's enough to fix one source s , and compute the min-cut between s and every other node.

proof:

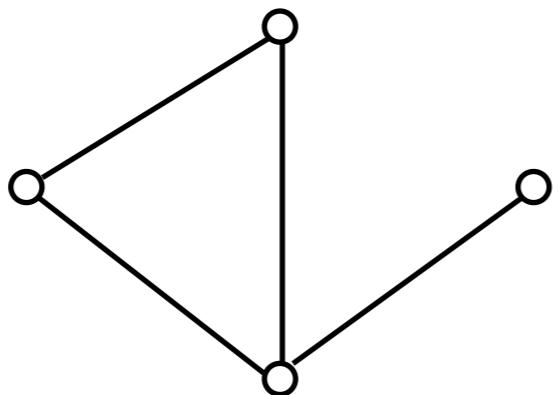
- in any min-cut A and $V-A$, the node s is assigned to one of the two. wlog we can assume that s is assigned to A .
- let v be some other node in A
- since s and v are both in A , this means that any cut separating the two is larger than the min-cut
- This implies that if v were the source instead of s , the same set A would have been found.

Randomized algorithm for min-cut

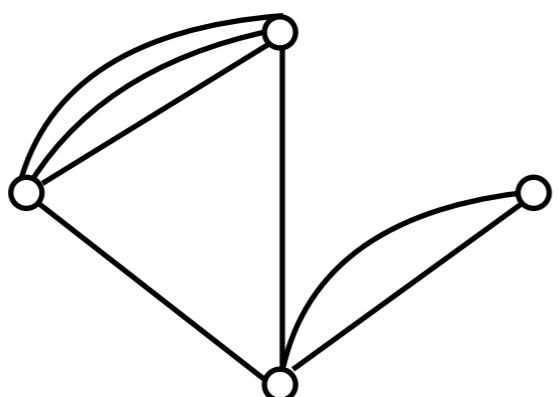
- $O(nmD)$ is polynomial, but slow
 - if $n = 10^6$, $m = O(n^2)$, $D = 10$, then we get $O(10^{19})$
- idea: make random choices
 - give up on finding an exact solution, i.e. the true min-cut
 - do we really give up...?
 - make it very fast
 - we can get $O(n)$
- result: the algorithm will find the min-cut with high(ish) probability
 - trick: repeat the algorithm multiple times and return the best of the outputs

multi graphs

simple graph: each pair of vertices are connected by at most one edge



multi graph: vertices may have multiple edges between them

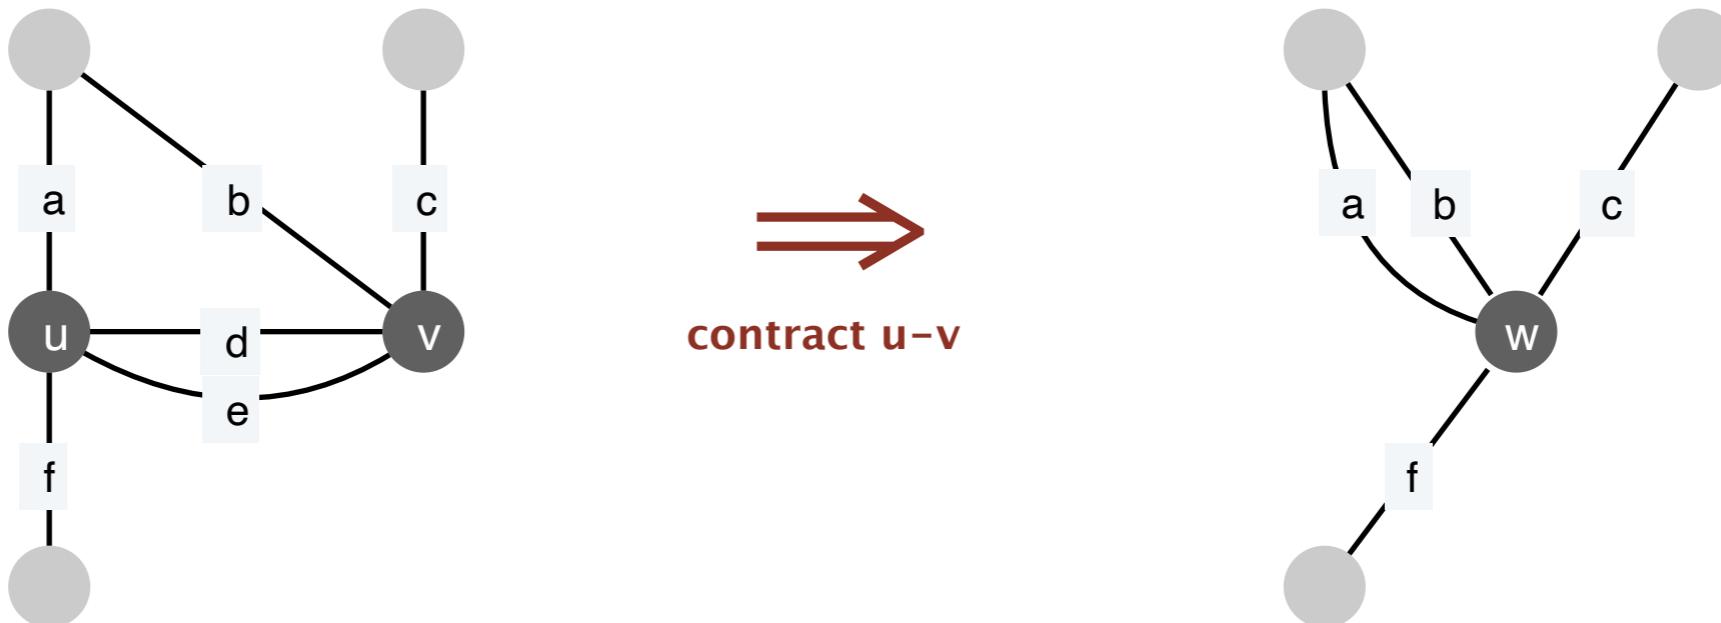


Contraction algorithm

Contraction algorithm. [Karger 1995]

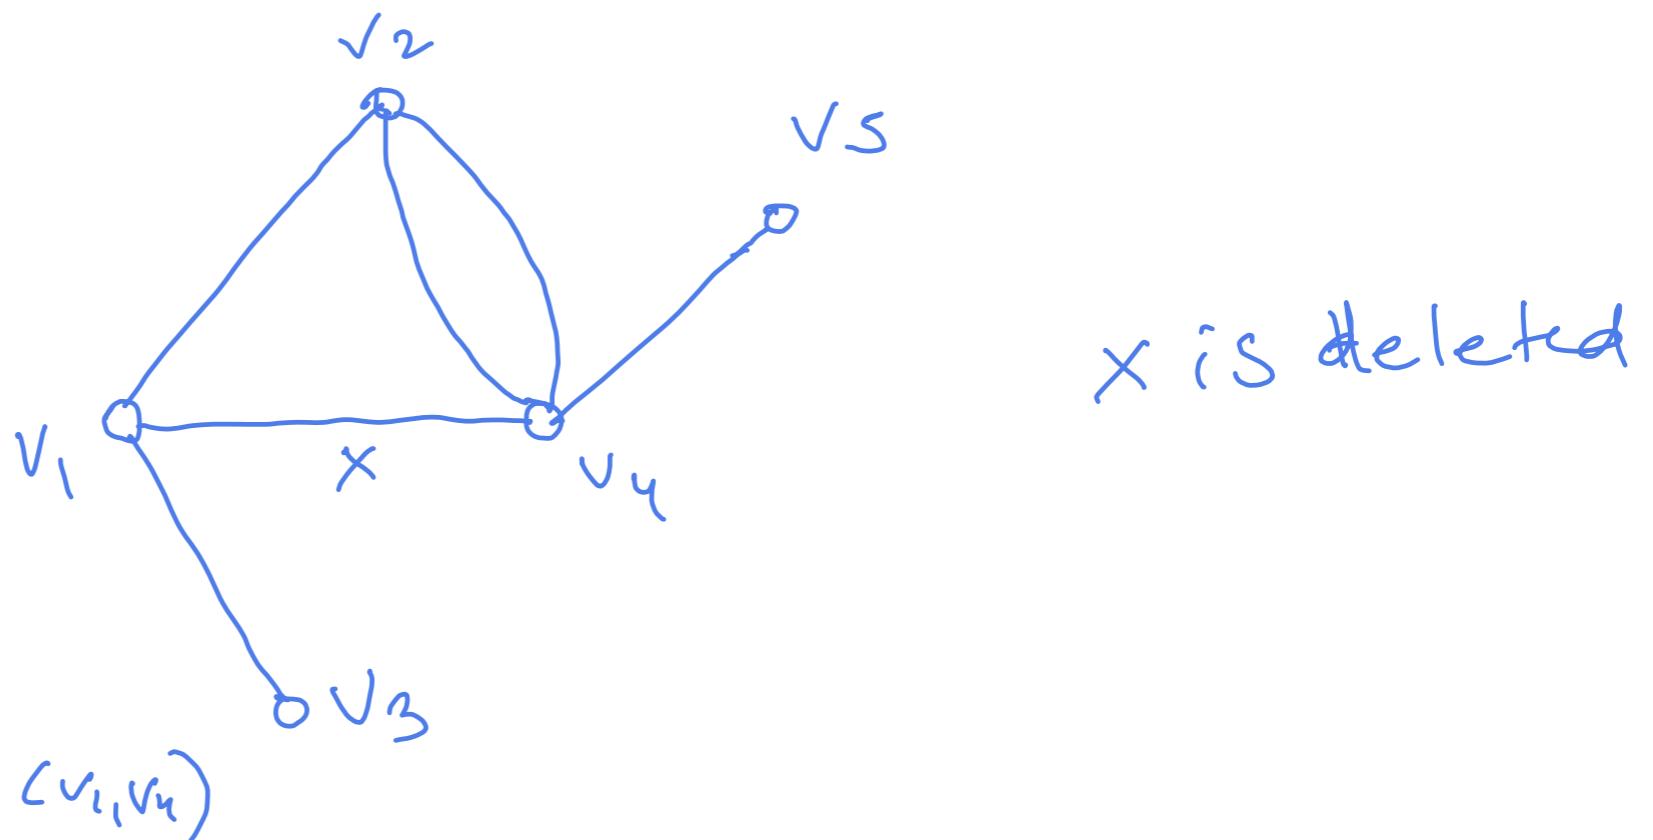
contraction operation:

- select an edge (u,v) at random
- replace u and v by a single super-node w
 - hereby deleting all edges between (u,v)
 - all edges previously adjacent to either u or v are now connected to w
 - keep parallel edges but delete self loops

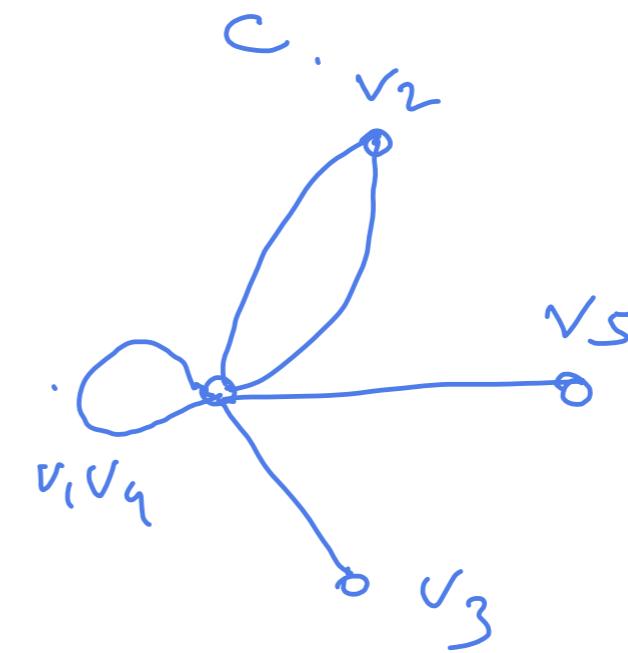
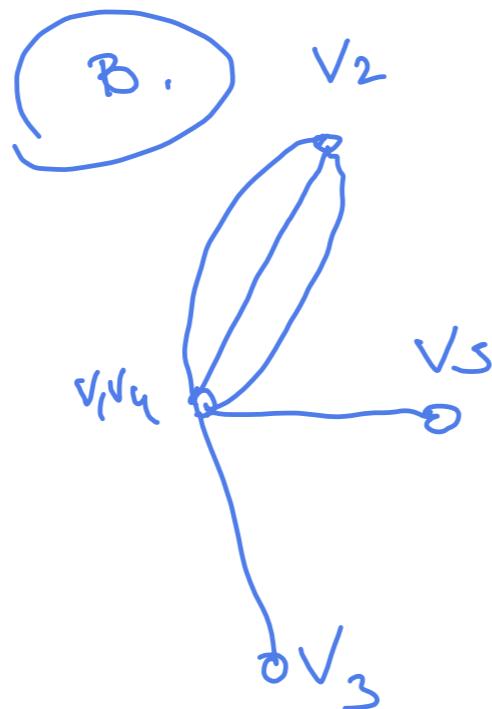
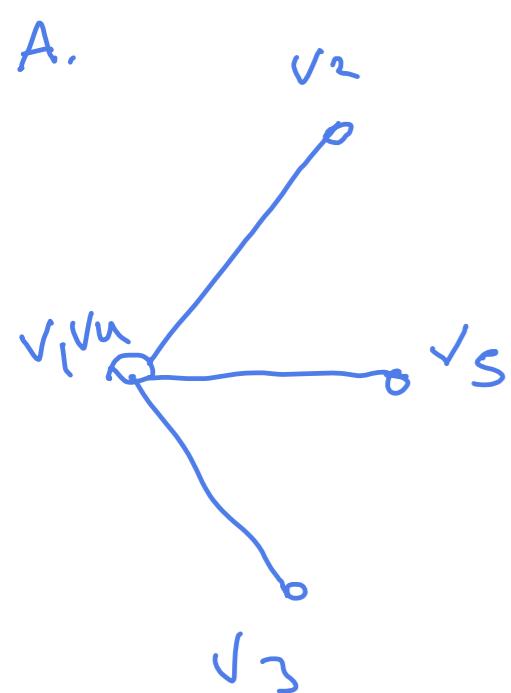


TopHat - vertex contraction

Here is a graph:



Suppose we contract edge $(\underline{u}, \underline{v})$, what is the corresponding multigraph?

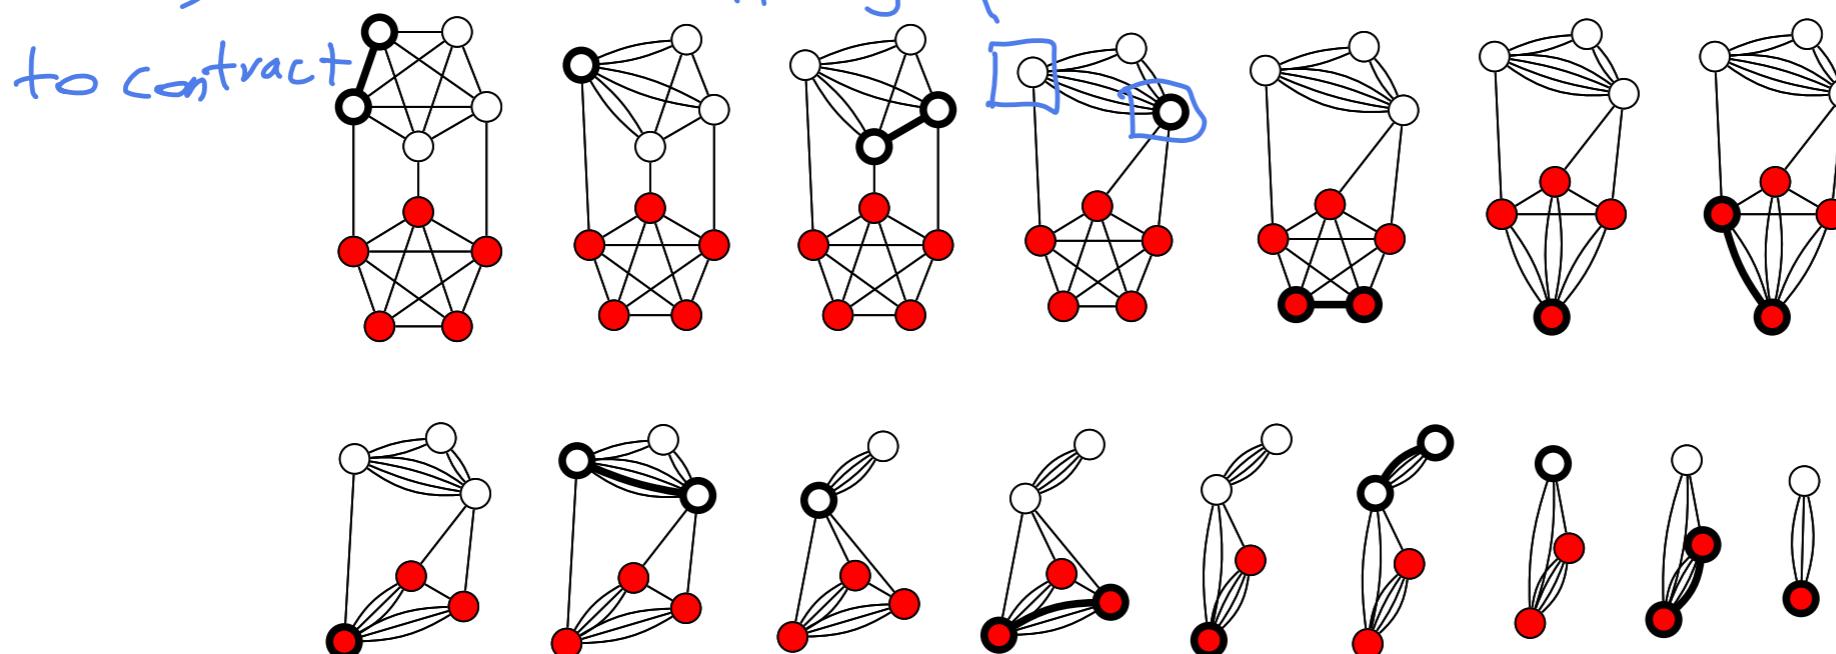


Contraction algorithm

Contraction algorithm. [Karger 1995]

- Pick an edge $e = (u, v)$ uniformly at random.
- Contract edge e .
 - replace u and v by single new super-node w
 - preserve edges, updating endpoints of u and v to w
 - keep parallel edges, but delete self-loops
- Repeat until graph has just two nodes v_1 and v_2 .
- Return the cut (all nodes that were contracted to form v_1).

every contraction
results in a
graph with
one fewer
nodes



Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

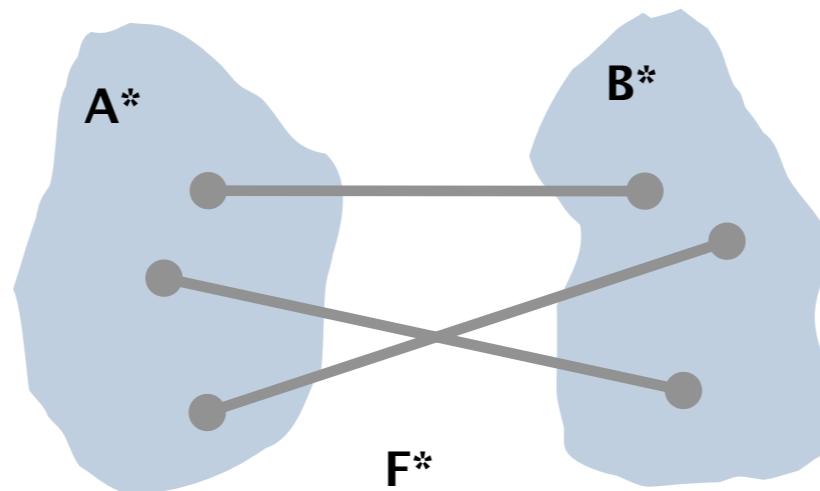
the contraction algorithm finds a min-cut, iff none of the edges in the min-cut get contracted.

goal: compute the probability of the event that the min-cut edges are not selected for contraction

Intuition:

vertices connected by multi-edges are more likely to be contracted

implies that larger cuts are more likely to be contracted



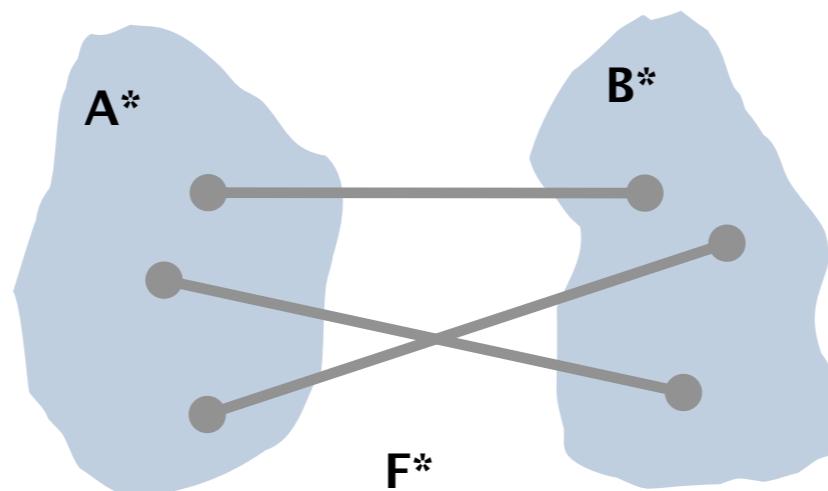
Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- Let F^* be edges with one endpoint in A^* and the other in B^* .
- Let $k = |F^*| = \text{size of min cut}$.
- pick a random edge to contract —> what is the probability that this edge is in F^* ?

the algo finds the min-cut if we never pick an edge in F^*



Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

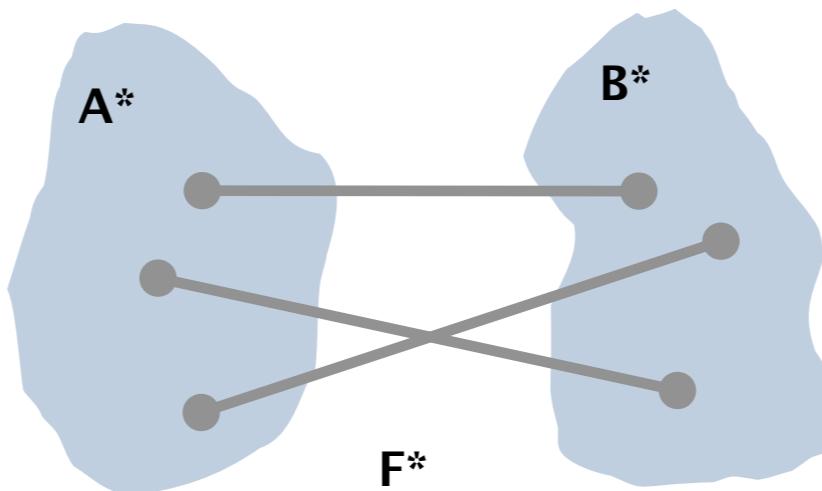
Pf. Consider a global min-cut (A^*, B^*) of G .

- F^* = edges in the min cut, $k = |F^*|$ = size of min cut.

- observation: each node has degree at least k

pf: k is the size of the min-cut
 $\rightarrow \deg(v) < k$ then separating v would
be a smaller cut.

- number of edges $|E| \geq \frac{k \cdot n}{2}$

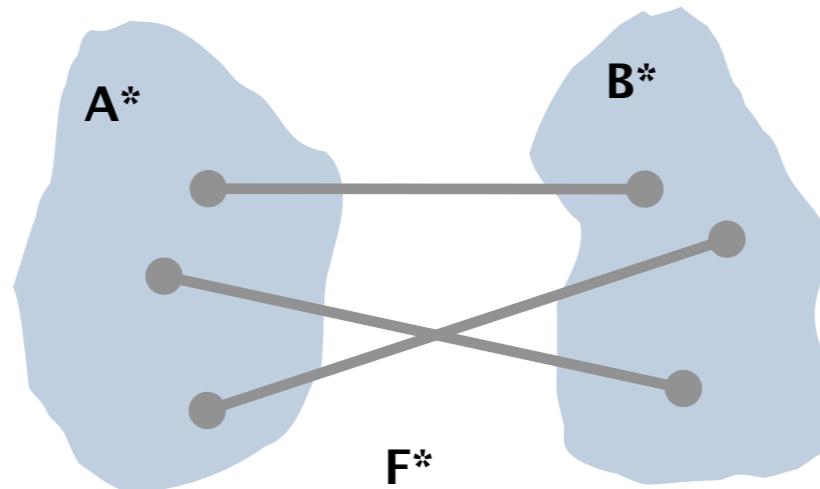


Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- $F^* = \text{edges in the min cut}, \quad k = |F^*| = \text{size of min cut}.$
- observation: each node has degree at least k
 - if some node v had degree less than k , then separating v from the rest of the graph creates a cut smaller than k
- number of edges $|E| \leq \frac{1}{2}kn$
 - each vertex has degree $\geq k$, there are n vertices. This way each edge is counted twice.



TopHat - probability of contracting F^*

If $|F^*| = k$, then we know that each node has degree $\geq k$, what is the probability of randomly picking an edge (u,v) in the global min-cut F^* ?

A. $P((u,v) \text{ in } F^*) = \frac{1}{n}$

$$P((u,v) \text{ in } F^*) = \frac{k}{|E|} \leq \frac{k}{\frac{1}{2}kn} = \frac{2}{n}$$

B. $P((u,v) \text{ in } F^*) = \frac{2}{n}$

C. $P((u,v) \text{ in } F^*) = \frac{k}{n}$

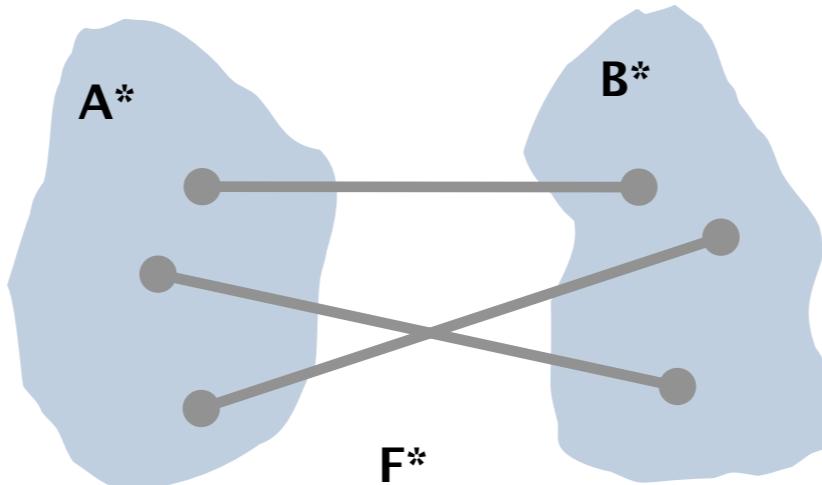
D. $P((u,v) \text{ in } F^*) = \frac{k}{n^2}$

Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- F^* = edges in the min cut, $k = |F^*|$ = size of min cut.
- In first step, algorithm contracts an edge in F^* probability $k / |E|$.
 - size of $|E|$?
 - Every node has degree $\geq k \Rightarrow |E| \geq \frac{1}{2} k n$.
- Thus, algorithm contracts an edge in F^* with probability $\leq \frac{k}{|E|} \stackrel{\textcolor{blue}{L}}{=} \frac{k}{\frac{1}{2} kn} = \frac{2}{n}$

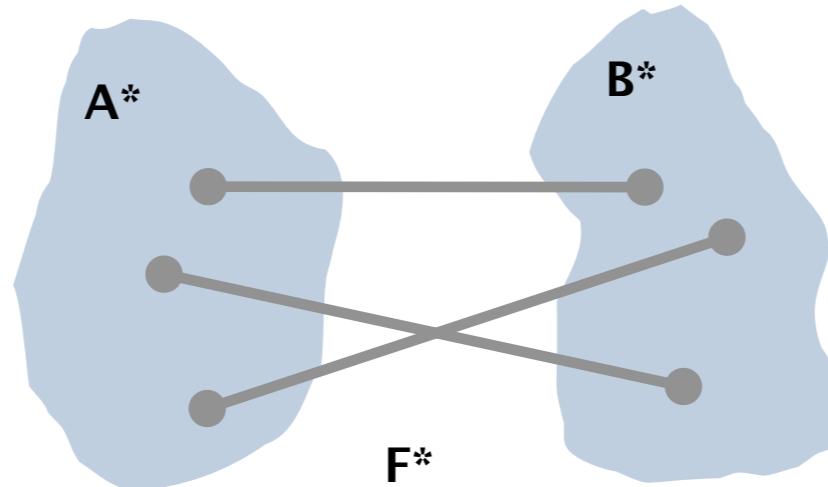


Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- $F^* = \text{edges in the min cut}, \quad k = |F^*| = \text{size of min cut}.$
- After j iterations, we have (contracted) graph G' with $n' = n-j$ nodes
- if none of the edges in F^* have been contracted:
 - the min-cut in G' still has size k
 - number of edges $|E'| \geq \frac{1}{2}kn'$
 - the algorithm contracts an edge in F^* with probability $\frac{2}{n'}$



Contraction algorithm – probability analysis

$$P(A|B) = \frac{P(B|A)P(A)}{P(B)}$$

Claim. The contraction algorithm returns a min cut with prob $\geq 2/n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- F^* = edges in the min cut, $k = |F^*|$ = size of min cut.
- After j iterations, we have (contracted) graph G' with $n' = n-j$ nodes
 - the algorithm contracts an edge in F^* with probability $\frac{2}{n'}$
 - E_j = event that an edge in F^* is *not* contracted in iteration j

$$P(E_1 \cap E_2 \cap \dots \cap E_{n-2}) = P(E_1) P(E_2 | E_1) P(E_3 | E_1 \wedge E_2) \dots P(E_{n-2} | E_1 \wedge \dots \wedge E_{n-3})$$

chain rule

all events have to happen so that no edge in F^* gets contracted

$$\begin{aligned} & \Rightarrow \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \left(1 - \frac{2}{n-2}\right) \dots \left(1 - \frac{2}{3}\right) = \\ & = \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \left(\frac{n-4}{n-2}\right) \left(\frac{n-5}{n-3}\right) \dots \frac{2}{4} \cdot \frac{1}{3} = \frac{2}{n(n-1)} \geq \frac{2}{n^2} \end{aligned}$$

Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

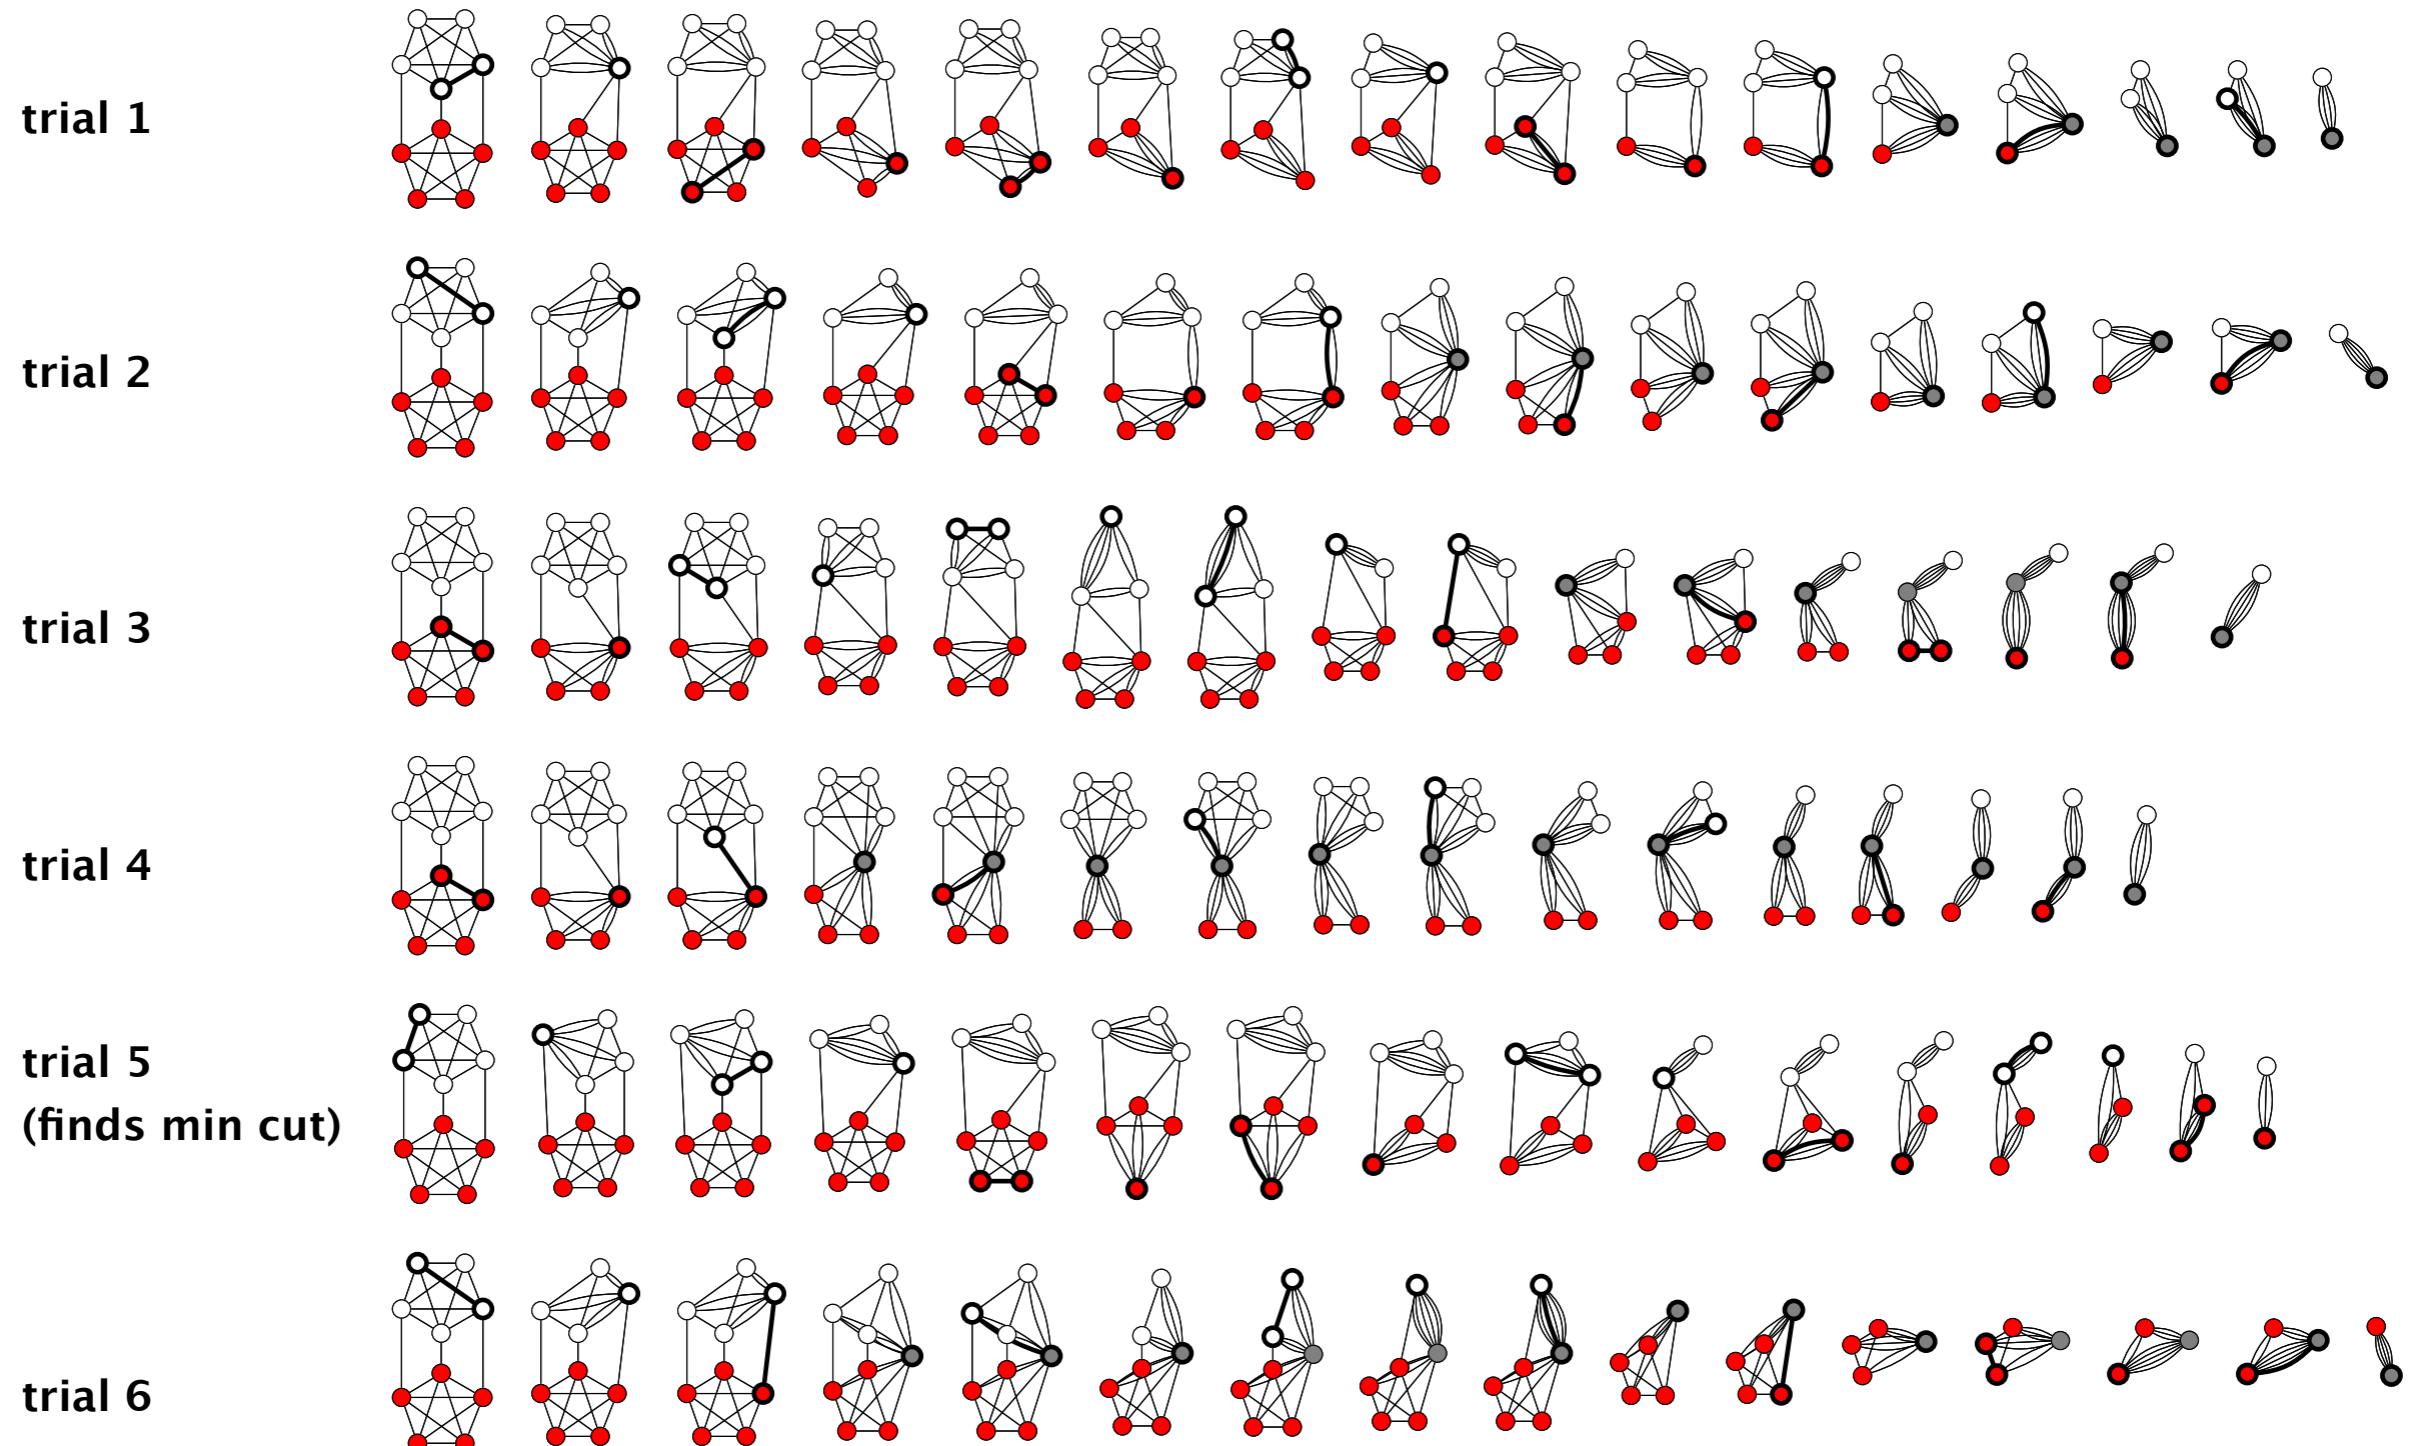
Pf. Consider a global min-cut (A^*, B^*) of G .

- $F^* = \text{edges in the min cut}, k = |F^*| = \text{size of min cut}.$
- After j iterations, we have (contracted) graph G' with $n' = n-j$ nodes
 - the algorithm contracts an edge in F^* with probability $\frac{2}{n'}$
 - $E_j = \text{event that an edge in } F^* \text{ is } not \text{ contracted in iteration } j$

$$\begin{aligned} P(E_1 \cap E_2 \cap \dots \cap E_{n-2}) &= P(E_1) \cdot P(E_2 | E_1) \cdot \dots \cdot P(E_{n-2} | E_1 \cap E_2 \cap \dots \cap E_{n-3}) \\ &\stackrel{\text{chain rule}}{\geq} \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \dots \left(1 - \frac{2}{4}\right) \left(1 - \frac{2}{3}\right) \\ &\stackrel{\text{substitute } 2/n'}{=} \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \left(\frac{n-4}{n-2}\right) \dots \left(\frac{2}{4}\right) \left(\frac{1}{3}\right) \\ &= \frac{2}{n(n-1)} \geq \frac{2}{n^2} \end{aligned}$$

Contraction algorithm: example execution

Amplification. To amplify the probability of success, run the contraction algorithm many times.



Reference: Thore Husfeldt

Contraction algorithm

Amplification. To amplify the probability of success, run the contraction algorithm many times.

Claim. If we repeat the contraction algorithm $n^2 \ln n$ times, then the probability of failing to find the global min-cut is $\leq 1/n^2$.

with independent random choices,

$$\left(1 - \frac{1}{x}\right)^x \leq \frac{1}{e}$$

Pf. By independence, the probability of failure is at most

$$\left(1 - \frac{2}{n^2}\right)^{n^2 \ln n} = \left[\left(1 - \frac{2}{n^2}\right)^{\frac{1}{2}n^2}\right]^{2 \ln n} \leq \left(e^{-1}\right)^{2 \ln n} = \frac{1}{n^2} \quad \begin{matrix} \uparrow \\ (1 - 1/x)^x \leq 1/e \end{matrix} \quad \begin{matrix} \text{tiny if} \\ n \text{ is large} \end{matrix}$$

In practice :

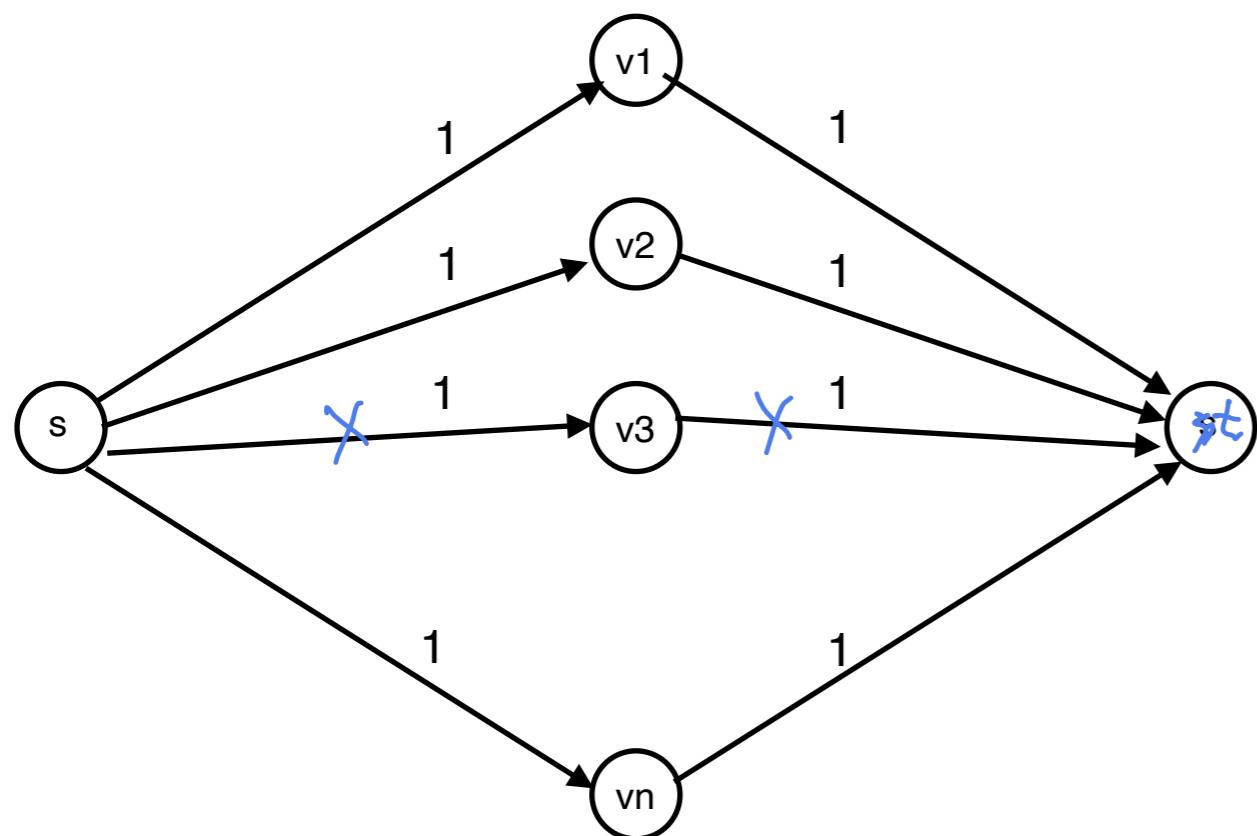
The first \sqrt{n} iterations are likely to be lucky
→ run that many to get some progress
⇒ then use different method
⇒ clever combination of this algo

TopHat - number of min st-cuts in directed graph

How many min st-cuts does this graph have?

- A. n B. $n \log n$ C. n^2

D. 2^n



n path from s to t



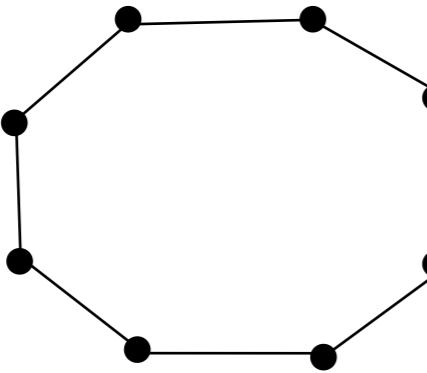
→ each such path
we have two
choices to cut:
either (s, v_i)
or (v_i, t)

$$\underbrace{2 \cdot 2 \cdot 2 \cdots 2}_{n \text{ times}} = 2^n$$

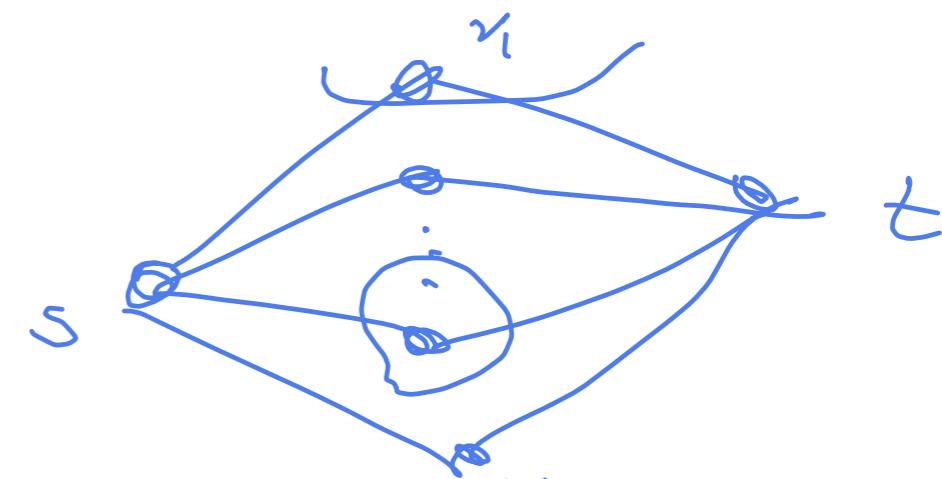
Exponential!

max number of global min-cuts in undirected graphs

How many global min-cuts on a cycle of length- n ? \equiv



one min-cut = cut any pair of edges
 $\Rightarrow \binom{n}{2} = \frac{n(n-1)}{2} = O(n^2)$



Separating one middle node is a min-cut
 $\Rightarrow n$ min-cuts!

How many global min-cuts (in worst case) are there in an undirected unweighted graph?

- guesss:
- 2^n
 - n^2
 - $n \log n$
 - n

Union bound

Event space (universe) Ω of all possible outcomes of a random trial

event $E_i \subseteq \Omega$ one (set of) possible outcome

example of events:

- roll 2 with a dice
- roll an even number with a dice
- process p_i is unsuccessful at accessing the database after t attempts
- the

intuitively: union of sets is less than the sum of individual sets

Union bound: given events E_1, E_2, \dots, E_r we have $P\left(\bigcup_{i=1}^r E_i\right) \leq \sum_{i=1}^r P(E_i)$

Number of global min-cuts

Claim: An undirected graph on n nodes has at most $\binom{n}{2}$ global min-cuts. = the cycle example!

pf: G is an undirected, unweighted graph

- C_1, C_2, \dots, C_r r min-cuts in G .
- E_i = event that C_i happens to be returned by the contraction algo

• $\Pr(E_i) \geq \frac{2}{n^2} \leftarrow$

• we don't care which min-cut it returns:
need union of events E_i

$$1 \geq \Pr\left(\bigcup_{i=1}^r E_i\right) \stackrel{\text{union bound}}{\geq} \sum_{i=1}^r \Pr(E_i) \geq r \cdot \frac{2}{n^2} \geq \frac{r}{n^2}$$

because

Prob.

$$1 \geq \frac{r}{n^2} \Rightarrow n^2 \geq r$$

• union bound
• we can use = because
the events of returning
any of the cuts is indep.

Number of global min-cuts

Claim: An undirected graph on n nodes has at most $\binom{n}{2}$ global min-cuts.

proof:

- C_1, C_2, \dots, C_r are the global min-cuts
- Let E_i be the event that C_i is returned by the Contraction Algorithm
- And $E = \bigcup_{i=1}^r E_i$ the event that the contraction algorithm returns a min-cut
- we know $P(E_i) \geq \frac{2}{n^2}$ from previous proof
- $P(E) = P(\bigcup_{i=1}^r E_i) = \sum_{i=1}^r P(E_i) \geq \frac{r}{n^2}$ Union bound suggests \leq , however the events E_i are disjoint (as the contraction algorithm can only return one of C_i in one specific run)
- finally $1 \geq P(E) = \frac{r}{n^2} \rightarrow n^2 \geq r$

Global min cut: context

Remark. Overall running time is slow since we perform $\Theta(n^2 \log n)$ iterations and each takes $\Omega(m)$ time.

Improvement. [Karger–Stein 1996] $O(n^2 \log^3 n)$.

- Early iterations are less risky than later ones: probability of contracting an edge in min cut hits 50% when $n / \sqrt{2}$ nodes remain.
- Run contraction algorithm until $n / \sqrt{2}$ nodes remain.
- Run contraction algorithm **twice** on resulting graph and return **best** of two cuts.

Extensions. Naturally generalizes to handle positive weights.

Best known. [Karger 2000] $O(m \log^3 n)$. ←

faster than best known max flow algorithm or
deterministic global min cut algorithm

Best known deterministic. [Nagamochi-Ibaraki 1992] $O(mn + n^2 \log n)$