

CS630 Graduate Algorithms

October 22, 2024

by Dora Erdos and Jeffrey Considine

- min-cut KT ch. 13.2

min st-cut in directed graphs

Max Flow Min Cut theorem: given a directed graph $G(V,E)$ with source s , sink t and edge capacities the value of the max flow is = the capacity of the min st-cut

\Rightarrow find the min-cut through the max-flow.

min st-cut in undirected unweighted graph

Find the min st-cut in the undirected, unweighted graph $G(V,E)$ with source s and sink t .

capacity/weight of a cut = number of edges between S and $V-S$.

How?

Global min-cut

Global min-cut: Given an undirected graph $G(V,E)$, find a cut that partitions the nodes into two sets A and $V-A$ with minimum weight.

- no designated source or sink

$$\bullet \text{ } weight(C) = \sum_{\text{edge } e \text{ in } C} w(e)$$

- works for unweighted and weighted graphs (today: unweighted)

Deterministic algorithm (using st-cuts):

Global min-cut

Global min-cut: Given an undirected graph $G(V,E)$, find a cut that partitions the nodes into two sets A and $V-A$ with minimum weight.

- no designated source or sink

$$\text{weight}(C) = \sum_{\text{edge } e \text{ in } C} w(e)$$

- works for unweighted and weighted graphs (today: unweighted)

Deterministic algorithm:

For each pair of vertices u, v assign u as source v as sink, and run the min st-cut algorithm

- since this is an undirected graph, it doesn't matter which of u and v is the source

running time:

- there are $O(n^2)$ iterations
- each requires a run of FF. We know that C is at most the max degree D .
- in total $O(n^2mD)$

Global min-cut

Global min-cut: Given an undirected graph $G(V,E)$, find a cut that partitions the nodes into two sets A and $V-A$ with minimum weight.

Deterministic algorithm — speed up:

claim: it's enough to fix one source s , and compute the min-cut between s and every other node.

runtime:

Global min-cut

Global min-cut: Given an undirected graph $G(V,E)$, find a cut that partitions the nodes into two sets A and $V-A$ with minimum weight.

Deterministic algorithm — speed up:

claim: it's enough to fix one source s , and compute the min-cut between s and every other node.

proof:


- in any min-cut A and $V-A$, the node s is assigned to one of the two. wlog we can assume that s is assigned to A .
- let v be some other node in A
- since s and v are both in A , this means that any cut separating the two is larger than the min-cut
- This implies that if v were the source instead of s , the same set A would have been found.

Randomized algorithm for min-cut


- $O(nmD)$ is polynomial, but slow
 - if $n = 10^6$, $m = O(n^2)$, $D = 10$, then we get $O(10^{19})$
- idea: make random choices
 - give up on finding an exact solution, i.e. the true min-cut
 - do we really give up...?
 - make it very fast
 - we can get $O(n)$
- result: the algorithm will find the min-cut with high(ish) probability
 - trick: repeat the algorithm multiple times and return the best of the outputs

multi graphs

simple graph: each pair of vertices are connected by at most one edge



multi graph: vertices may have multiple edges between them




Contraction algorithm


Contraction algorithm. [Karger 1995]

contraction operation:

- select an edge (u,v) at random
- replace u and v by a single super-node w
 - hereby deleting all edges between (u,v)
 - all edges previously adjacent to either u or v are now connected to w
 - keep parallel edges but delete self loops





contract $u-v$



Contraction algorithm

Contraction algorithm. [Karger 1995]

- Pick an edge $e = (u, v)$ uniformly at random.
- **Contract** edge e .
 - replace u and v by single new super-node w
 - preserve edges, updating endpoints of u and v to w
 - keep parallel edges, but delete self-loops
- Repeat until graph has just two nodes v_1 and v_2 .
- Return the cut (all nodes that were contracted to form v_1).



Reference: Thore Husfeldt

Contraction algorithm – probability analysis


Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

the contraction algorithm finds a min-cut, iff none of the edges in the min-cut get contracted.

goal: compute the probability of the event that the min-cut edges are not selected for contraction

Intuition:

vertices connected by multi-edges are more likely to be contracted
implies that larger cuts are more likely to be contracted




Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- Let F^* be edges with one endpoint in A^* and the other in B^* .
- Let $k = |F^*|$ = size of min cut.
- pick a random edge to contract \rightarrow what is the probability that this edge is in F^* ?




Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- F^* = edges in the min cut, $k = |F^*|$ = size of min cut.
- observation: each node has degree at least k
- number of edges $|E| =$




Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- F^* = edges in the min cut, $k = |F^*|$ = size of min cut.
- observation: each node has degree at least k
 - if some node v had degree less than k , then separating v from the rest of the graph creates a cut smaller than k
- number of edges $|E| = \frac{1}{2}kn$
 - each vertex has degree $\geq k$, there are n vertices. This way each edge is counted twice.




Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- F^* = edges in the min cut, $k = |F^*|$ = size of min cut.
- In **first step**, algorithm contracts an edge in F^* probability $k / |E|$.
 - size of $|E|$?
 - Every node has degree $\geq k \Rightarrow |E| \geq \frac{1}{2}kn$
- Thus, algorithm contracts an edge in F^* with probability \leq




Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- F^* = edges in the min cut, $k = |F^*|$ = size of min cut.
- After j iterations, we have (contracted) graph G' with $n' = n - j$ nodes
- if none of the edges in F^* have been contracted:
 - the min-cut in G' still has size k
 - number of edges $|E'| \geq \frac{1}{2}kn'$
 - the algorithm contracts an edge in F^* with probability $\frac{2}{n'}$



Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

Pf. Consider a global min-cut (A^*, B^*) of G .

- F^* = edges in the min cut, $k = |F^*|$ = size of min cut.
- After j iterations, we have (contracted) graph G' with $n' = n - j$ nodes
 - the algorithm contracts an edge in F^* with probability $\frac{2}{n'}$
 - E_j = event that an edge in F^* is *not* contracted in iteration j

$$P(E_1 \cap E_2 \cap \dots E_{n-2}) =$$

Contraction algorithm – probability analysis

Claim. The contraction algorithm returns a min cut with prob $\geq 2 / n^2$.

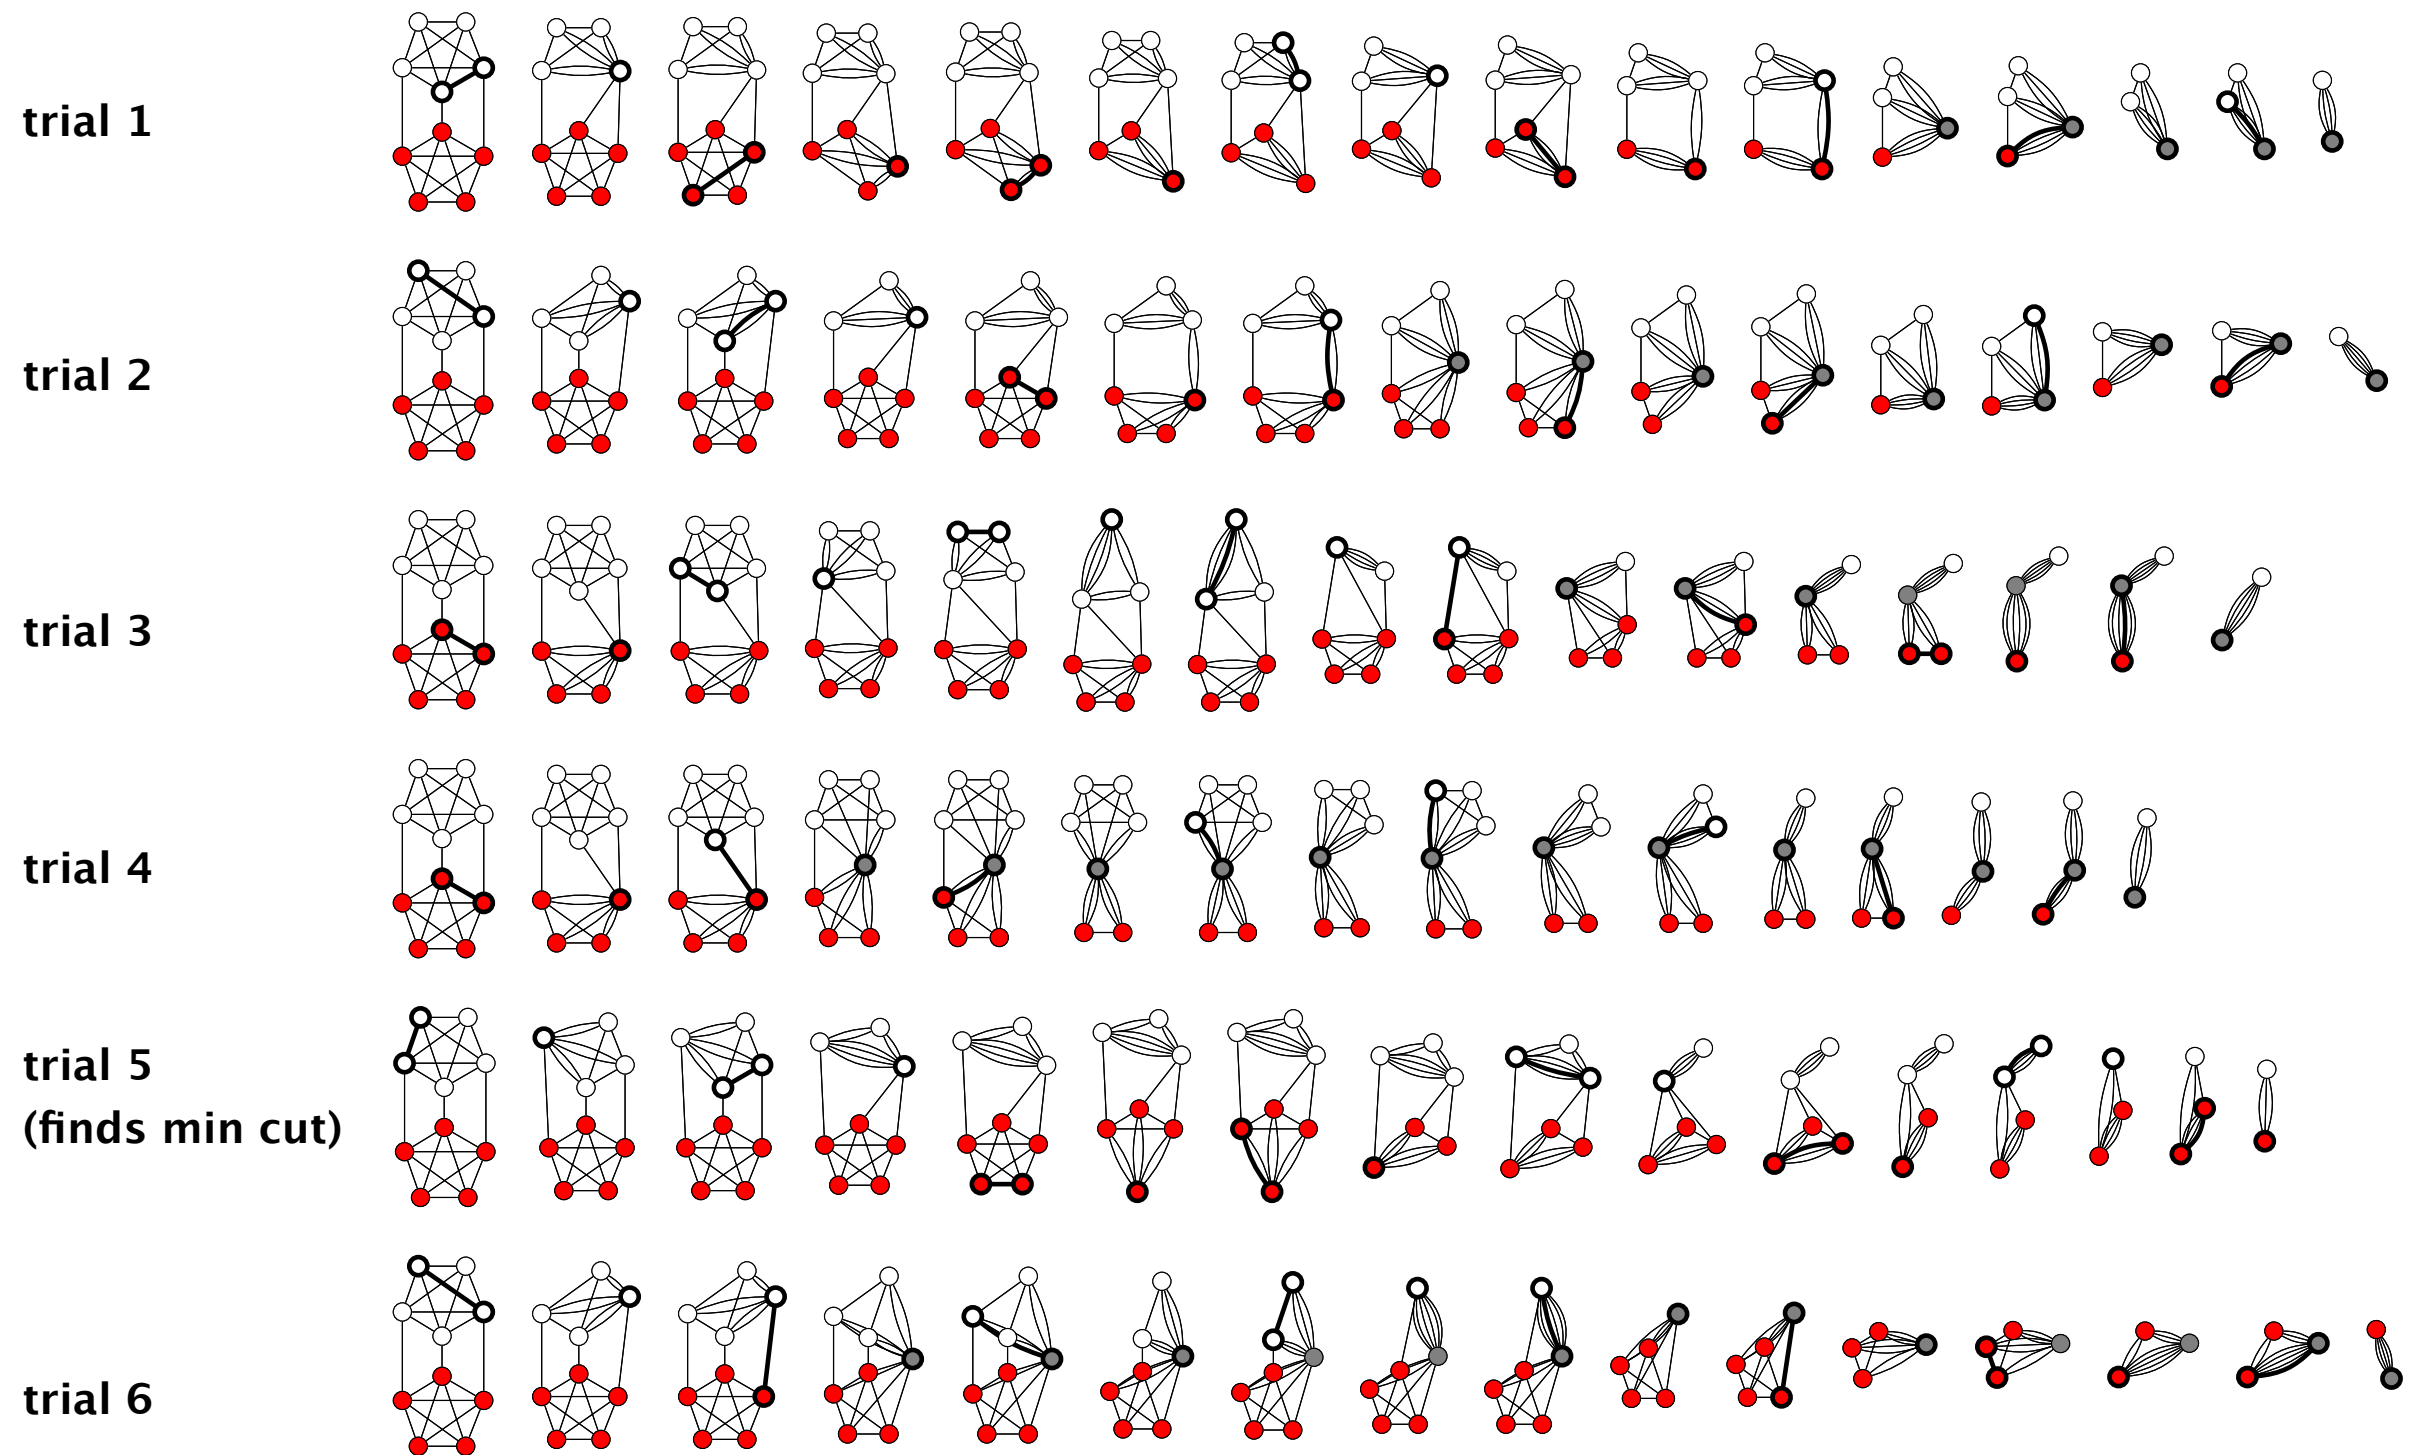
Pf. Consider a global min-cut (A^*, B^*) of G .

- F^* = edges in the min cut, $k = |F^*|$ = size of min cut.
- After j iterations, we have (contracted) graph G' with $n' = n - j$ nodes
 - the algorithm contracts an edge in F^* with probability $\frac{2}{n'}$
 - E_j = event that an edge in F^* is *not* contracted in iteration j

$$\begin{aligned} P(E_1 \cap E_2 \cap \dots E_{n-2}) &= P(E_1) \cdot P(E_2 | E_1) \cdot \dots \cdot P(E_{n-2} | E_1 \cap E_2 \cap \dots \cap E_{n-3}) \\ &\stackrel{\text{chain rule}}{\geq} \left(1 - \frac{2}{n}\right) \left(1 - \frac{2}{n-1}\right) \dots \left(1 - \frac{2}{4}\right) \left(1 - \frac{2}{3}\right) \\ &\stackrel{\text{substitute } 2/n'}{=} \left(\frac{n-2}{n}\right) \left(\frac{n-3}{n-1}\right) \left(\frac{n-4}{n-2}\right) \dots \left(\frac{2}{4}\right) \left(\frac{1}{3}\right) \\ &= \frac{2}{n(n-1)} \geq \frac{2}{n^2} \end{aligned}$$

Contraction algorithm: example execution

Amplification. To amplify the probability of success, run the contraction algorithm many times.



Reference: Thore Husfeldt

Contraction algorithm

Amplification. To amplify the probability of success, run the contraction algorithm many times.

Claim. If we repeat the contraction algorithm $n^2 \ln n$ times, then the probability of failing to find the global min-cut is $\leq 1 / n^2$.

with independent random choices,




Pf. By independence, the probability of failure is at most

$$\left(1 - \frac{2}{n^2}\right)^{n^2 \ln n} = \left[\left(1 - \frac{2}{n^2}\right)^{\frac{1}{2}n^2}\right]^{2 \ln n} \leq \left(e^{-1}\right)^{2 \ln n} = \frac{1}{n^2}$$

\uparrow
 $(1 - 1/x)^x \leq 1/e$

max number of global min-cuts in undirected graphs

How many global min-cuts on a cycle of length- n ?



How many global min-cuts (in worst case) are there in an undirected unweighted graph?

Union bound

Event space (universe) Ω of all possible outcomes of a random trial

event $E_i \subseteq \Omega$ one (set of) possible outcome

example of events:

- roll 2 with a dice
- roll an even number with a dice
- process p_i is unsuccessful at accessing the database after t attempts
- the

intuitively: union of sets is less than the sum of individual sets

Union bound: given events E_1, E_2, \dots, E_r we have $P\left(\bigcup_{i=1}^r E_i\right) \leq \sum_{i=1}^r P(E_i)$

Number of global min-cuts

Claim: An undirected graph on n nodes has at most $\binom{n}{2}$ global min-cuts.


Number of global min-cuts

Claim: An undirected graph on n nodes has at most $\binom{n}{2}$ global min-cuts.

proof:

- C_1, C_2, \dots, C_r are the global min-cuts
- Let E_i be the event that C_i is returned by the Contraction Algorithm
- And $E = \cup_{i=1}^r E_i$ the event that the contraction algorithm returns *a* min-cut

- we know $P(E_i) \geq \frac{2}{n^2}$ from previous proof

- $P(E) = P(\cup_{i=1}^r E_i) = \sum_{i=1}^r P(E_i) \geq \frac{r}{n^2}$  Union bound suggests \leq , however the events E_i are independent (as the contraction algorithm can only return one of C_i in one specific run)
- finally $1 \geq P(E) = \frac{r}{n^2} \rightarrow n^2 \geq r$


Global min cut: context

Remark. Overall running time is slow since we perform $\Theta(n^2 \log n)$ iterations and each takes $\Omega(m)$ time.

Improvement. [Karger–Stein 1996] $O(n^2 \log^3 n)$.

- Early iterations are less risky than later ones: probability of contracting an edge in min cut hits 50% when $n / \sqrt{2}$ nodes remain.
- Run contraction algorithm until $n / \sqrt{2}$ nodes remain.
- Run contraction algorithm **twice** on resulting graph and return **best** of two cuts.

Extensions. Naturally generalizes to handle positive weights.

Best known. [Karger 2000] $O(m \log^3 n)$.  faster than best known max flow algorithm or deterministic global min cut algorithm

Best known deterministic. [Nagamochi-Ibaraki 1992] $O(mn + n^2 \log n)$